

Contents

1	Machine Learning Foundations	1
1.1	SVD	1
1.1.1	Explanation	1
1.1.2	Construction of SVD	2
1.1.3	Implementation details	3

1 Machine Learning Foundations

1.1 SVD

1.1.1 Explanation

The singular value decomposition (SVD) is a matrix decomposition method which decomposes a matrix A into three specific matrices. The matrix A could contain complex numbers, but I will focus on the real numbers because the decomposition is used for machine learning.

Let $A \in \mathbb{R}^{m \times n}$, then there exists a decomposition $A = USV^T$ containing the following matrices:

- $U \in \mathbb{R}^{m \times m}$: an orthogonal matrix ($U^T U = U U^T = I$) with column vectors $\vec{u}_1, \dots, \vec{u}_m$.
- $V \in \mathbb{R}^{n \times n}$: an orthogonal matrix with column vectors $\vec{v}_1, \dots, \vec{v}_n$.
- $S \in \mathbb{R}^{m \times n}$: a rectangular matrix with only non-zero elements on the main diagonal, $S_{ii} \geq 0$ and $S_{ij} = 0, i \neq j$.

$$A = \begin{pmatrix} | & | & | & | \\ \vec{u}_1 & \vec{u}_2 & \dots & \vec{u}_m \\ | & | & | & | \end{pmatrix} \begin{pmatrix} \sigma_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_n \\ \vdots & & \vdots \\ 0 & \dots & 0 \end{pmatrix} \begin{pmatrix} | & | & | & | \\ \vec{v}_1 & \vec{v}_2 & \dots & \vec{v}_n \\ | & | & | & | \end{pmatrix}^T$$

The diagonal elements of S are called the **singular values**, noted as $S_{ii} = \sigma_i$. By convention, they are ordered as $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r \geq 0$, with r the rank. The columns of U are called the **left singular vectors**, and the columns of V are called the **right singular vectors**. If $m < n$, the matrix S would contain 'padding' of zeros on the right side, instead of underneath. It is needed because the matrix is the same size as the original A . SVD is a very generalized decomposition, so it **exists for any** $A \in \mathbb{R}^{m \times n}$.

1.1.2 Construction of SVD

The SVD is a generalized version of the diagonalization of square matrices. A square matrix $A_{n \times n}$ is called diagonalizable if there exists an invertible matrix $P_{n \times n}$ for which $A = PDP^{-1}$, with D a diagonal matrix containing the eigenvalues of A , and P containing the basis of eigenvectors as columns. Calculating the SVD of this $A_{n \times n}$ matrix will result in exactly the same decomposition as the diagonalization.

Computing the SVD of the rectangular $A_{m \times n}$ matrix is equivalent to finding two orthonormal bases $U = (\vec{u}_1, \dots, \vec{u}_m)$ and $V = (\vec{v}_1, \dots, \vec{v}_n)$. Start with constructing the right singular vectors $\vec{v}_1, \dots, \vec{v}_n$. Create the matrix $B := A^T A \in \mathbb{R}^{n \times n}$. B is symmetric because $B^T = (A^T A)^T = A^T (A^T)^T = A^T A = B$, and positive semi-definite, which means that for all $x : x^T B x \geq 0$. We obtain $x^T B x = x^T A^T A x = (Ax)^T (Ax) = \langle Ax, Ax \rangle$, which is positive because the dot product computes the sum of squares. Therefore, it is possible to diagonalize B as $B = V D V^T$, with D a diagonal matrix containing the eigenvalues $\theta_1, \dots, \theta_n$. Assume that the SVD exists such that:

$$B = A^T A = (U S V^T)^T (U S V^T) = V S^T U^T U S V^T = V S^T S V^T = V S^2 V^T$$

Remember that U is orthogonal so $U^T U = I$ and S is a diagonal matrix so $S^T = S$. We can now see that $B = V D V^T = V S^2 V^T$ thus $\theta_i = \sigma_i^2$. This concludes that V contains the right singular vectors.

To compute the left singular vectors, follow the same principle as above but with the matrix $C := A A^T \in \mathbb{R}^{m \times m}$:

$$C = A A^T = (U S V^T)(U S V^T)^T = U S V^T V S^T U^T = U S S^T U^T = U S^2 U^T$$

Because C can be diagonalized, we find $A A^T = U D U^T$ with U an orthonormal basis of eigenvectors, representing left singular vectors.

The matrix S can be constructed because we can take the square root of the eigenvalues from the previous diagonalization to find the singular values. The non-zero eigenvalues of $A A^T$ and $A^T A$ are the same, so the entries in S will also be the same.

To connect all the pieces, we have an orthonormal set of right singular vectors in the matrix V , the singular values in the matrix S , and an orthonormal set of left singular vectors in U . There is still a problem because these eigenvectors are not unique. If we first calculate V , normalisation is needed for U based on the vectors in V . This happens with the following calculations:

$$\vec{u}_i := \frac{A \vec{v}_i}{\|A \vec{v}_i\|} = \frac{1}{\sqrt{\theta_i}} A \vec{v}_i = \frac{1}{\sigma_i} A \vec{v}_i$$

This can be rearranged to find the formula:

$$A\vec{v}_i = \sigma_i u_i \quad \text{or} \quad AV = US$$

What results in the final SVD equation:

$$A = USV^T$$

1.1.3 Implementation details

The `transform` function takes a matrix A as input. This could be a matrix with any possible shape $m \times n$. The implementation first checks the shape of A for computational efficiency:

- When $m \geq n$, calculate $B := A^T A$.
- When $m < n$, calculate $B := AA^T$.

This way, we always work with the smallest matrix product to save computation time. To calculate the first singular vectors, we calculate the eigenvalue decomposition with the `eigh` function because it is specialized for symmetric/Hermitian matrices (which B always is). This is more numerically stable than the general `eig` solver. In theory, these eigenvalues only contain values ≥ 0 , but because of floating point arithmetic, small negative values could appear. These are rounded to zero if they occur. A similar problem arises with the singular values because some of them might be very small, resulting in near-zero division. A threshold is set to filter out that numerical noise. The singular values are calculated simply by taking the square root of the eigenvalues of B .

To compute the other singular vectors (depending on which case we are in), we use the formula $U = AV \frac{1}{\sigma}$ or $V = A^T U \frac{1}{\sigma}$, with $\frac{1}{\sigma}$ a diagonal matrix containing the normalized singular values. The function returns the three matrices U , S and V^T as final result.